

Distributed Systems Foundations

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Lecture 1


Main Characteristics of Distributed Systems

- Independent processors, sites, processes
- Message passing
- No shared memory
- No shared clock
- Independent failure modes

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



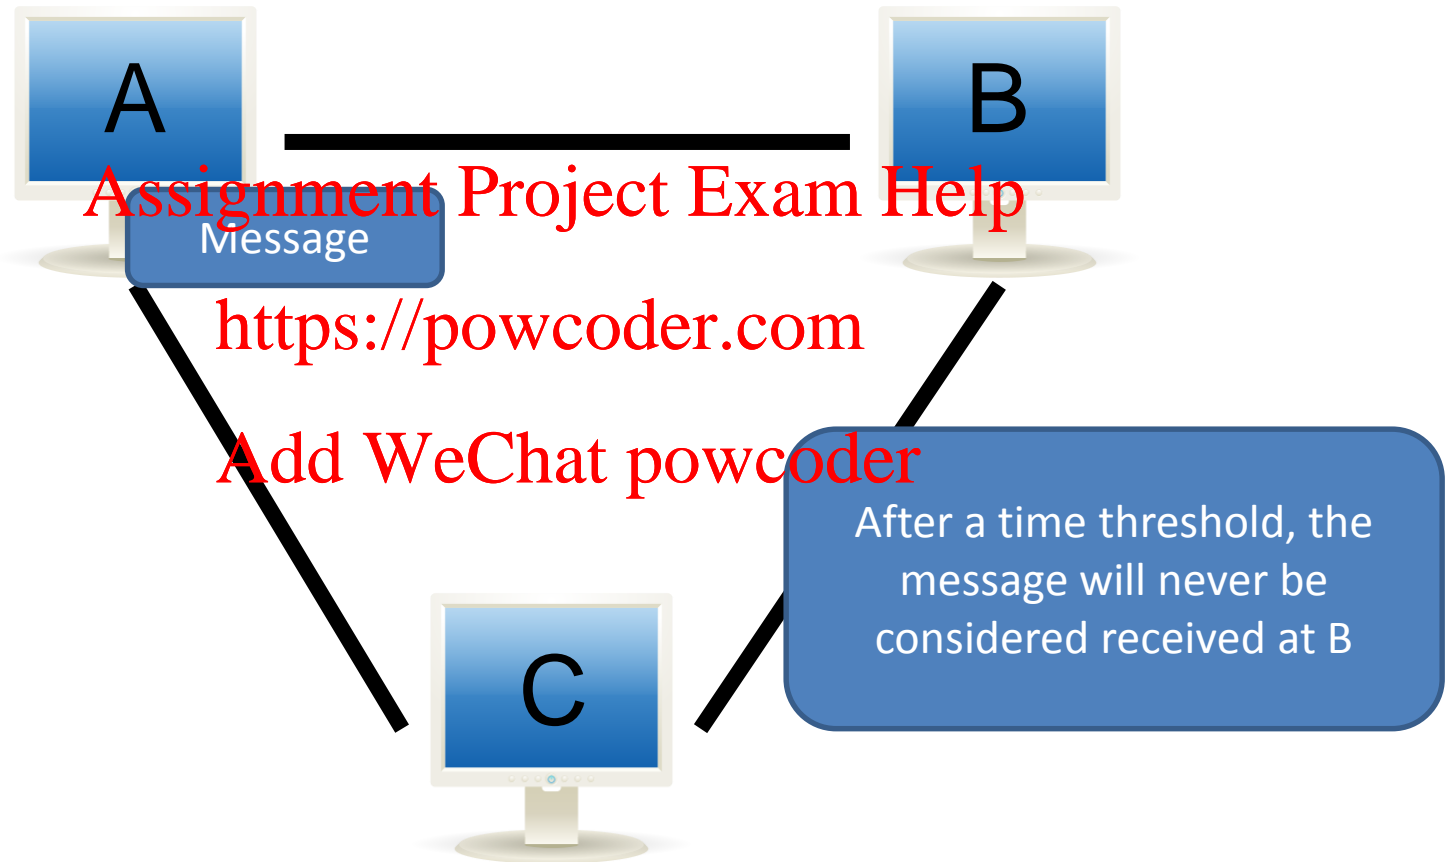
Distributed System Models

- **Synchronous System:** Known bounds on times for message transmission, processing, bounds on local clock drifts, etc.
– Can use timeouts <https://powcoder.com>
- **Asynchronous System:** No known bounds on times for message transmission, processing, bounds on local clock drifts, etc.
– More realistic, practical, but no timeout.

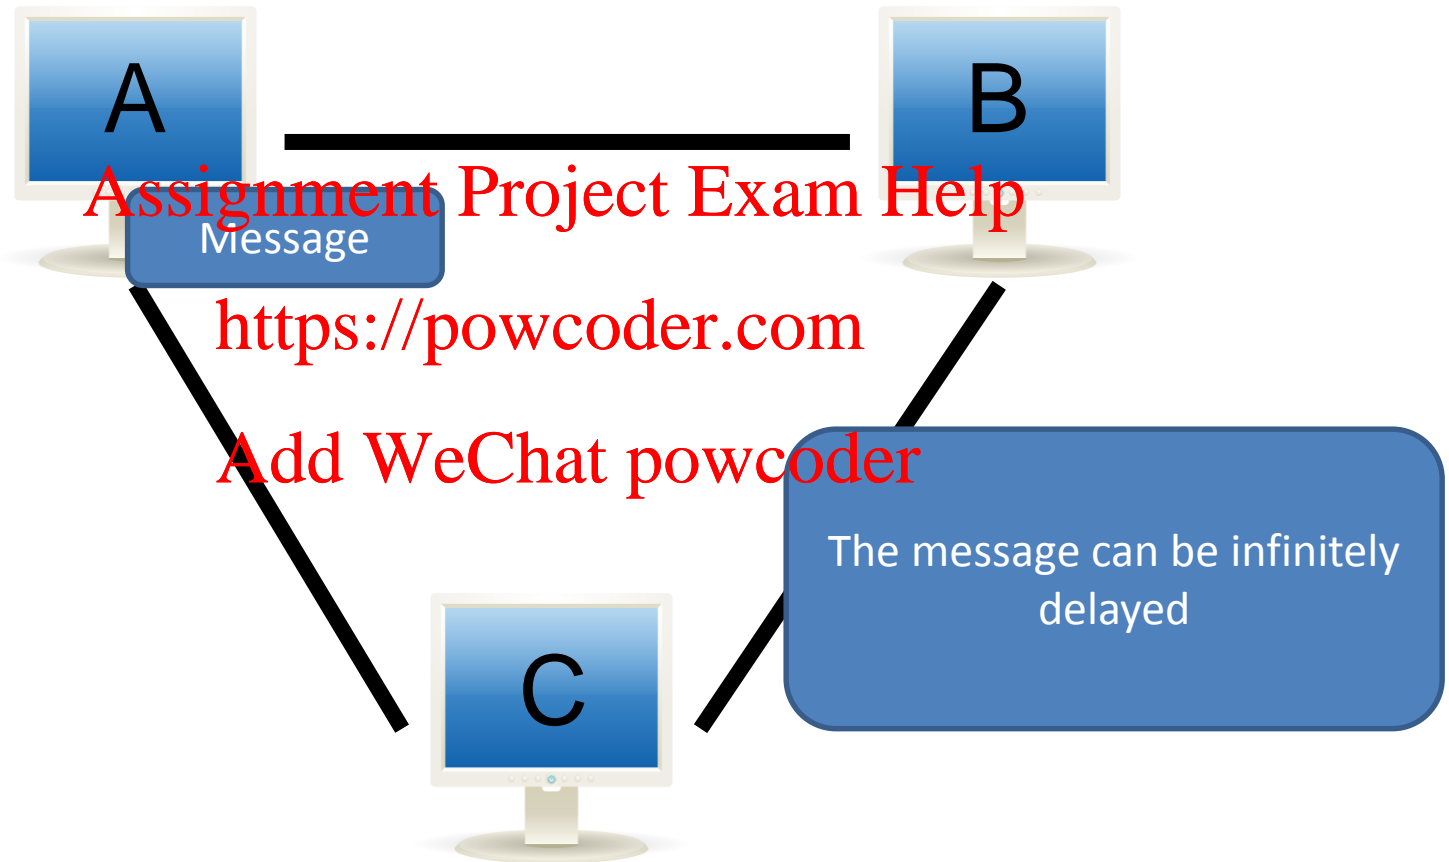
Distributed System Models

- **Synchronous System:** Known bounds on times for message transmission, processing, bounds on local clock drifts, etc.
– Can use timeouts <https://powcoder.com>
- **Asynchronous System:** No known bounds on times for message transmission, processing, bounds on local clock drifts, etc.
– More realistic, practical, but no timeout.


Synchronous model example



Asynchronous model example



Data link layer in a distributed system



Goal: Send a message from A to B

Data Link Layer Example



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- Assumptions:
 - Infinite buffer
 - No message loss
- Solution?

Data Link Layer Example



Assignment Project Exam Help

- Assumptions:

<https://powcoder.com>

- Finite buffer

- No message loss

Add WeChat powcoder

- Solution:

- Synchronous System?

- Asynchronous System?

Data Link Layer Example



Assignment Project Exam Help

- Assumptions:

<https://powcoder.com>

- Finite buffer

- Message loss

Add WeChat powcoder

- Solution?

Data Link Layer Lessons

**What are the
assumptions?**

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

CAUSALITY AND TIME

What is a Distributed System?

- A simple model of a distributed system proposed by **Lamport** in a landmark 1978 paper:
<https://powcoder.com/Assignment/Project/Exam/Help/>
- “Time, Clocks and the Ordering of Events in a Distributed System” Communications of the ACM

UCSB COMPUTER SCIENCE DEPARTMENT PRESENTS DISTINGUISHED LECTURE

LESLIE LAMPORT

/// MICROSOFT RESEARCH

**WHO
BUILDS
A
SKYSCRAPER
WITHOUT
DRAWING
BLUEPRINTS?**

Assignment Project Exam Help

<https://powcoder.com>

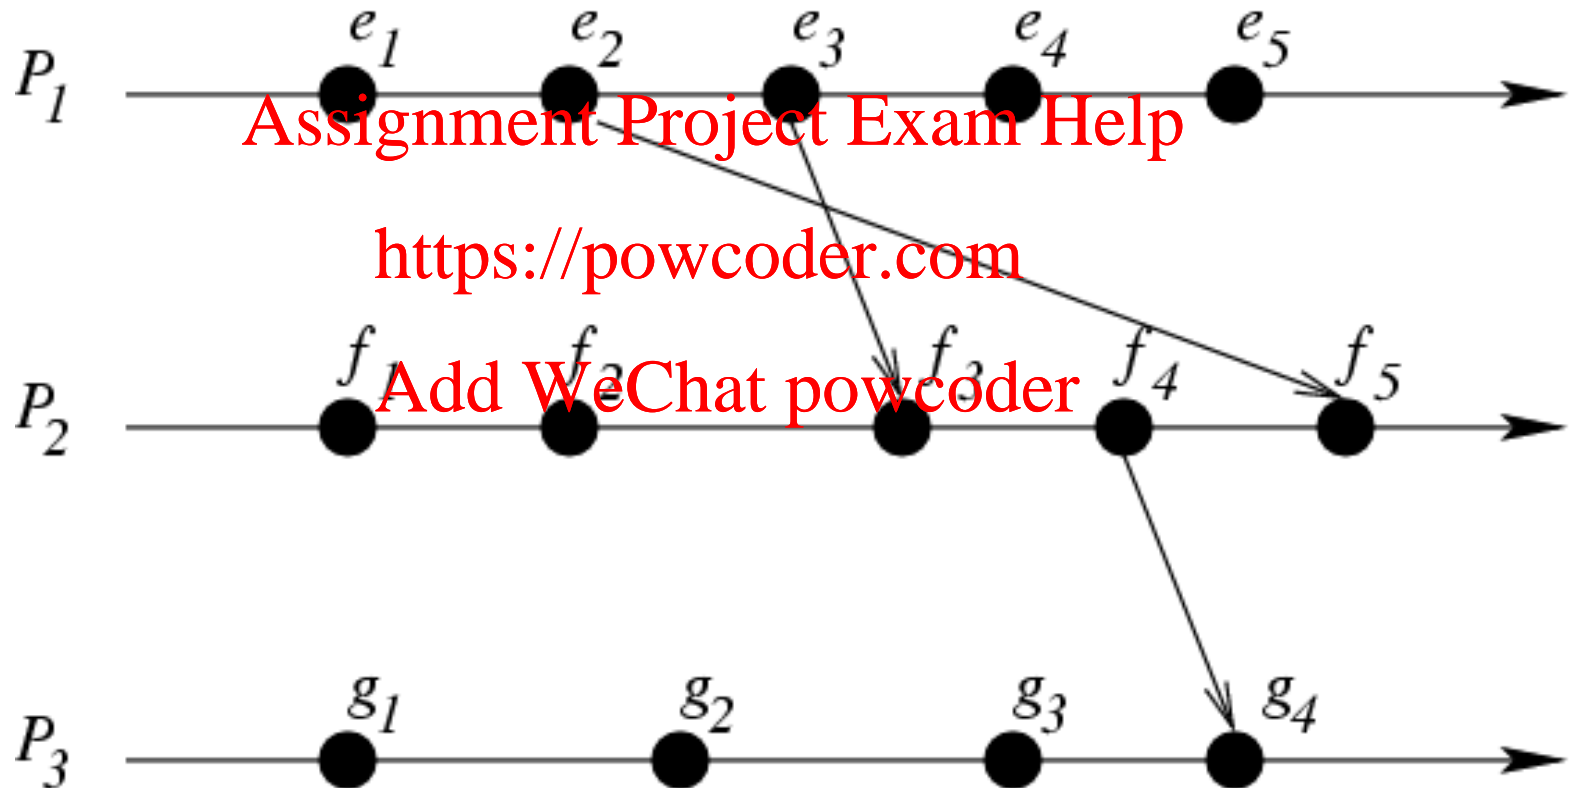
Add WeChat powcoder

 SPONSORED BY

What is a Distributed System?

- A set of processes that communicate using message passing.
- A process is a sequence of events
- 3 kinds of events:
 - Local events
 - Send events
 - Receive events
- Local events on a process for a total order.

Example of a Distributed System



Why do we care about “Time” in a distributed system?

- May need to know the time of day at which some event happens on a specific computer

- **external clock synchronization**

<https://powcoder.com>

- May need to know the relative order two events that happened on different computers

- **internal clock synchronization**

Physical Clocks in Distributed Systems

- Does this work?
 - Synchronize all the clocks to some known high degree of accuracy
 - Measure time relative to local clock to determine order between events

Assignment Project Exam Help

- Well, there are some problems...
 - It's difficult to synchronize the clocks
 - Crystal-based clocks tend to drift over time-count time at different rates, and diverge from each other
 - Physical variations in the crystals, temperature variations, etc.
 - Drift is small, but adds up over time
 - For quartz crystal time, typical drift rate is about one second every 10^6 seconds=11.6days
 - Best atomic clocks drift one second in 10^{13} seconds = 300,000 years

Logical Clocks

- Idea — abandon idea of physical time

Assignment Project Exam Help

- For many purposes, it is sufficient to know the **order** in which events occurred

<https://powcoder.com>

Add WeChat powcoder

- Lamport (1978) — introduce logical *time*, to provide **consistent event ordering**

ORDERING EVENTS

- Event ordering linked with **causality**:
 - Saying that event a happened before event b is same as saying that event a could have affected the outcome of event b
 - If events a and b happen on processes that do not exchange any data, their exact ordering is not important

Happens Before or Causal Order on Events

- Event e *happens before (causally precedes)* event f , denoted $e \rightarrow f$ if:
 - 1. The *same* process executes e before f ; or
 - 2. e is *send(m)* and f is *receive(m)*; or
 - 3. Exists h so that $e \rightarrow h$ and $h \rightarrow f$
- We define *concurrent*, $e \parallel f$, as:
 $\neg(e \rightarrow f \vee f \rightarrow e)$

Lamport Logical Clocks

- Assign “clock” value to each event such that
 - if $a \rightarrow b$ then $\text{clock}(a) < \text{clock}(b)$
- Assign each process a clock “counter”.
 - Clock is incremented between any two events in the same process
 - Each message carries the sender’s clock value
- When a message arrives set local clock to:
 - $\max(\text{local value}, \text{message timestamp}) + 1$
- This clock forms a partial order.

Logical Clocks

- Each event has a single integer as its logical clock
 - Each process has a local counter C
 - Increment C between two events
 - At each “send”, logical clock value V is attached.
 - At each “receive”, $C = \max(C, V) + 1$


Assignment Project Exam Help

<https://powcoder.com>

$$3 = \max(1, 2) + 1$$

$$4 = \max(3, 3) + 1$$


Add WeChat powcoder



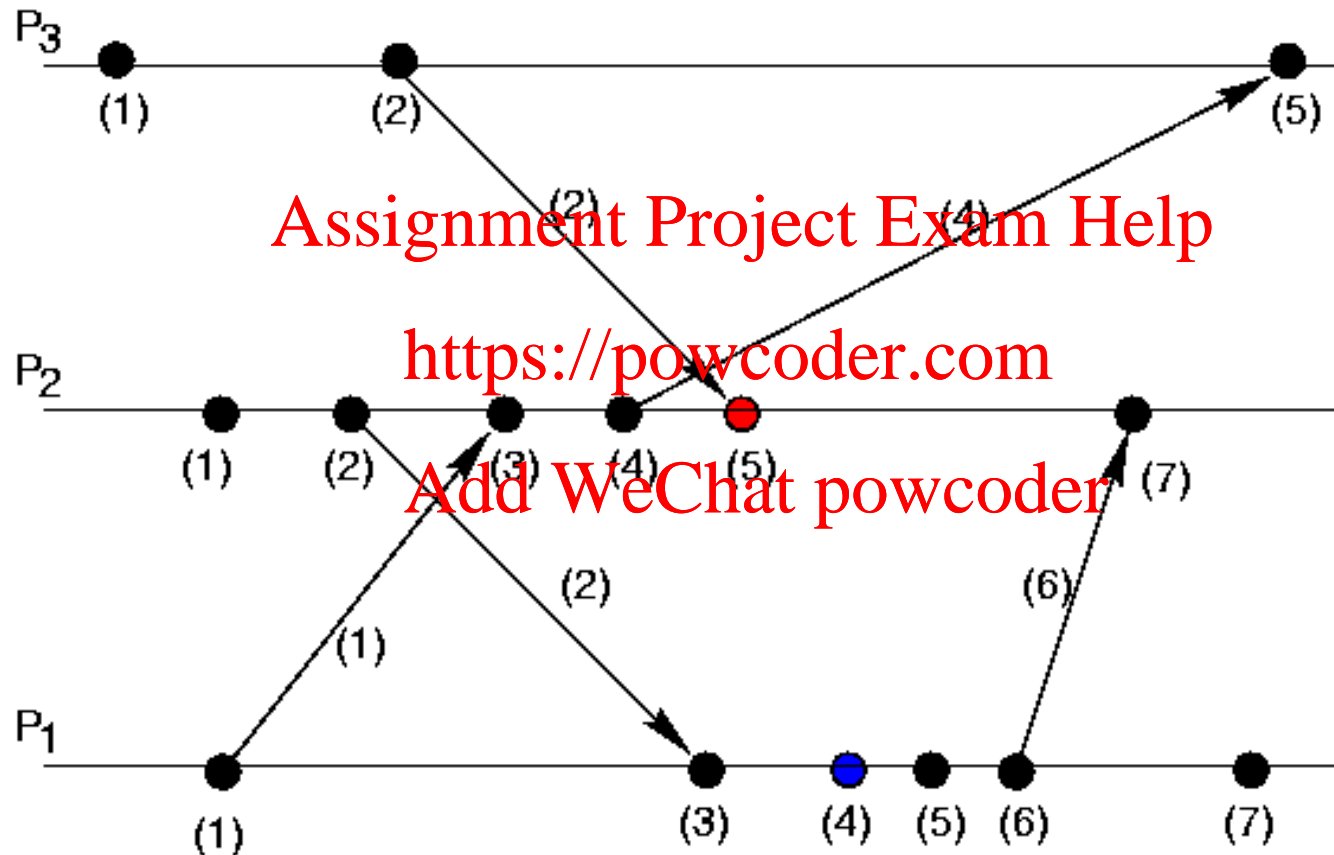
Logical Clock Properties

- Theorem:
 - Event s happens before $t \Rightarrow$ logical clock value of s is smaller than the logical clock value of t .

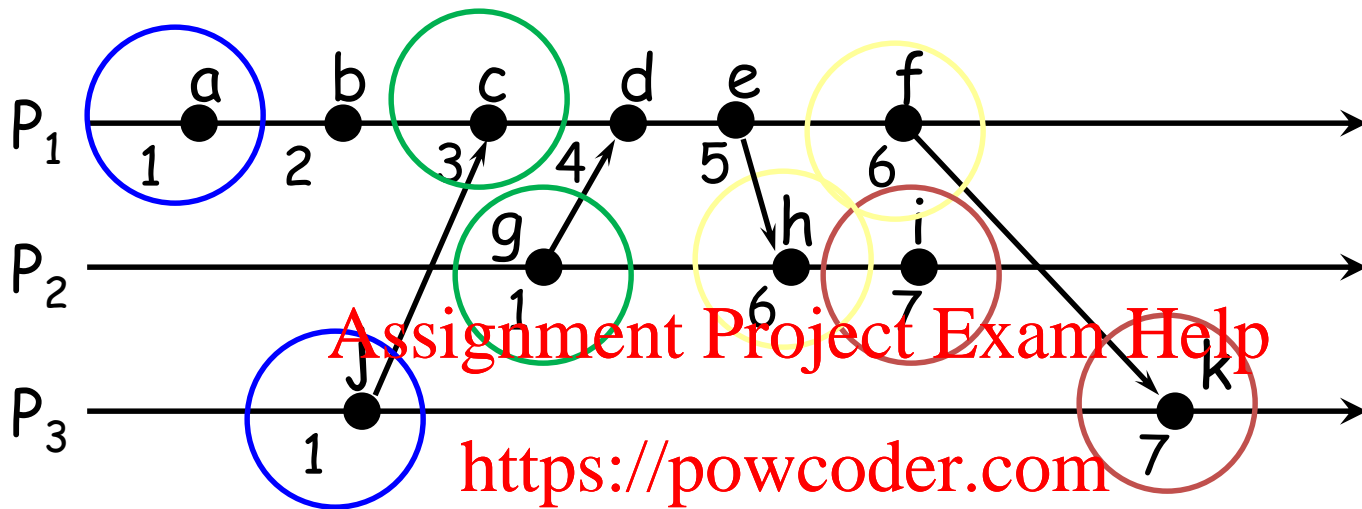
- The reverse may not be true



Example of a Logical Clock



Problem: Identical timestamps



$a \rightarrow b, b \rightarrow c, \dots$: Add WeChat events sequenced

$j \rightarrow c, g \rightarrow d, f \rightarrow k, \dots$: Lamport imposes a *send* \rightarrow *receive* relationship

Concurrent events (e.g., a & j) may have the same timestamp
 ... or not (e.g., c & g)


Total Order Lamport Clock

- To extend the partial order to total order: **use process ids to break ties**



- Global timestamps:
 - (T_a, P_a) : T_a is local Lamport logical timestamp
 P_a is process id.
 - $(T_a, P_a) < (T_b, P_b)$ iff
 - $(T_a < T_b)$ or $((T_a = T_b) \text{ and } (P_a < P_b))$
- Total order is consistent with partial order.
- Total order useful in many applications.

Unique (totally ordered) timestamps



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder