

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

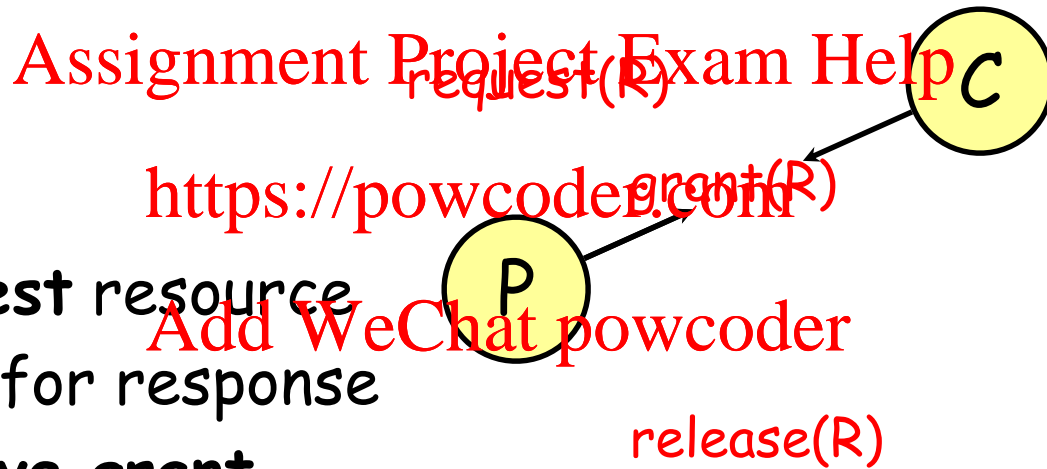
# MUTUAL EXCLUSION & QUORUMS

# Distributed Mutual Exclusion

- Given a set of processes and a single resource, develop a protocol to ensure exclusive access to the resource by a single process at a time.  
<https://powcoder.com>  
Add WeChat powcoder
- This is a fundamental operation in operating systems, and is generalized to locking in databases.

# Centralized algorithm

- One process elected as coordinator



1. **Request** resource
2. Wait for response
3. **Receive grant**
4. access resource
5. **Release** resource

# Centralized Solution

- But we can have multiple concurrent requests!
- Coordinator maintains queue of pending requests.
- Protocol:
  - Process send request to coordinator.
  - If no other request, coordinator sends back reply.
    - Otherwise, put request in queue
  - On receipt of reply, process accesses resource.
  - Once done, process sends release to coordinator.
  - On receipt of release, coordinator checks queue for any pending requests.

# Centralized Solution

Queue

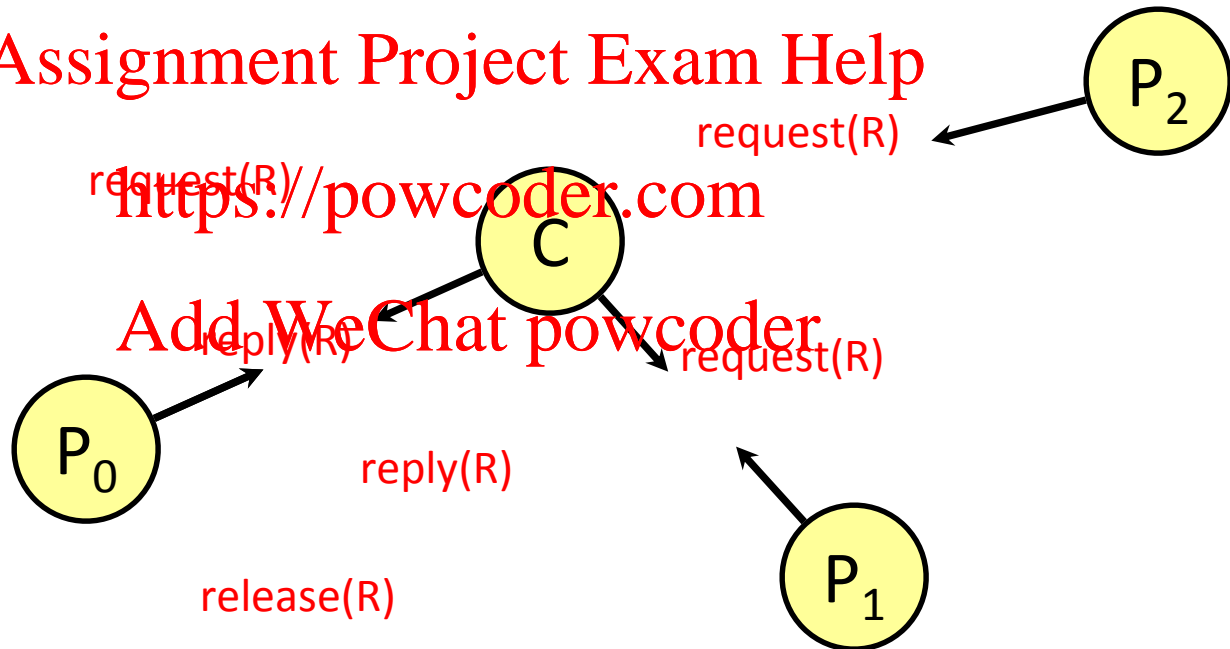
P<sub>1</sub>

P<sub>2</sub>

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



thanks paul krzyzanowski rutgers

# Centralized Approach

- Does it satisfy requests in order the requests are made?

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Lamport's Distributed Solution

- Instead of a central coordinator, **all processes collectively**
- Use **similar approach**, but engage **ALL** processes. <https://powcoder.com>
- Each request is time stamped with:
  - Lamport Logical Time + Identifier (process ID)
  - I.e., Timestamps are **totally-ordered Lamport** pairs.

# Lamport's Distributed Solution

- Process sends **time stamped request** to **all processes** and puts request in local queue.
- On receipt of request, **process** puts request in **queue** and sends back **reply**.
- Process **accesses resource**
  - On receipt of **reply**, process removes request at **head of queue**
  - Own request at **head of queue**
- Once done, process sends **release** to **all processes**.
- On receipt of release, **process** removes request



# Distributed Solution

- Does this work (Lamport original solution)?
- Need to order **queues** so they are **identical**:
  - Use **logical Lamport time** + **proc id** to break ties.
  - **FIFO** channels
- Requests are executed in **causal order**.
  - Why?


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

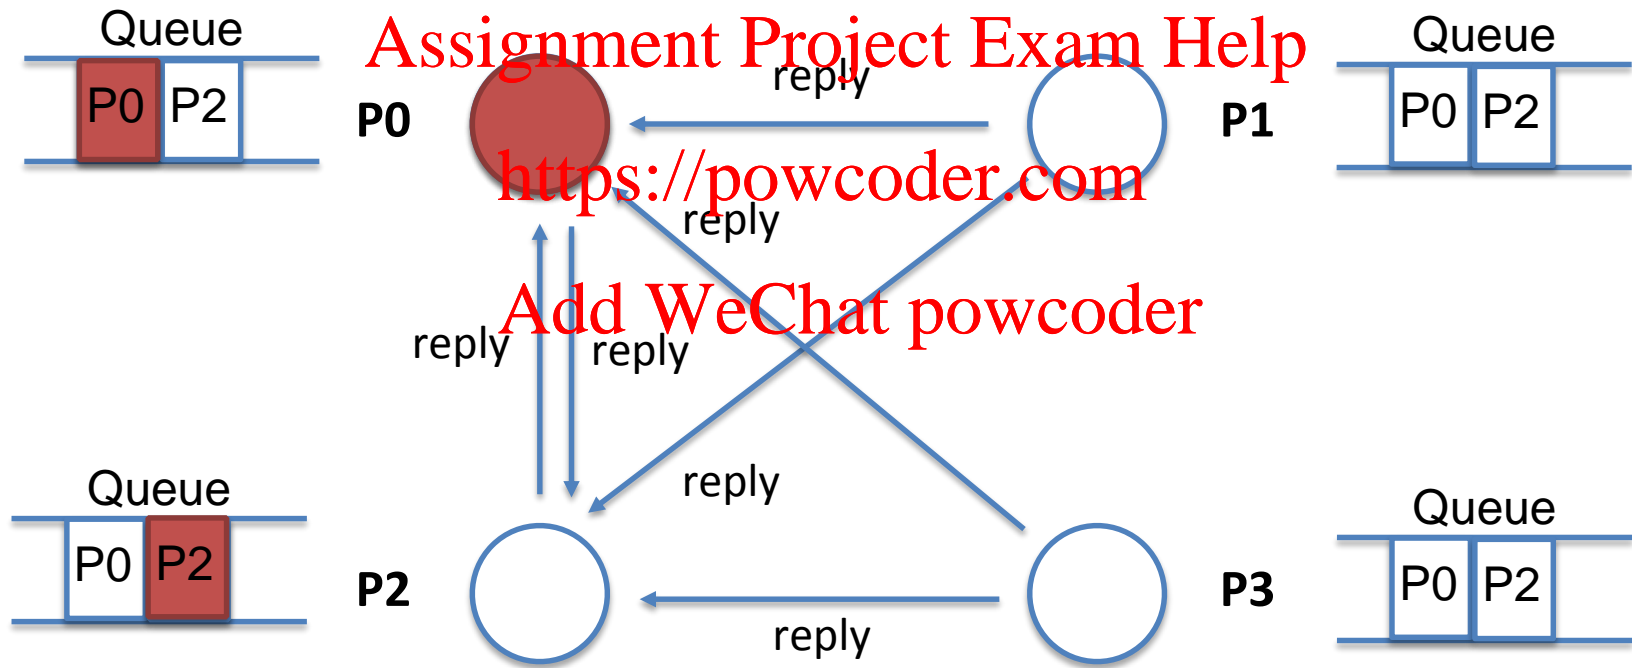
# Example

Assume that process P0 and process P2 need to access the shared resource at the **same time**.



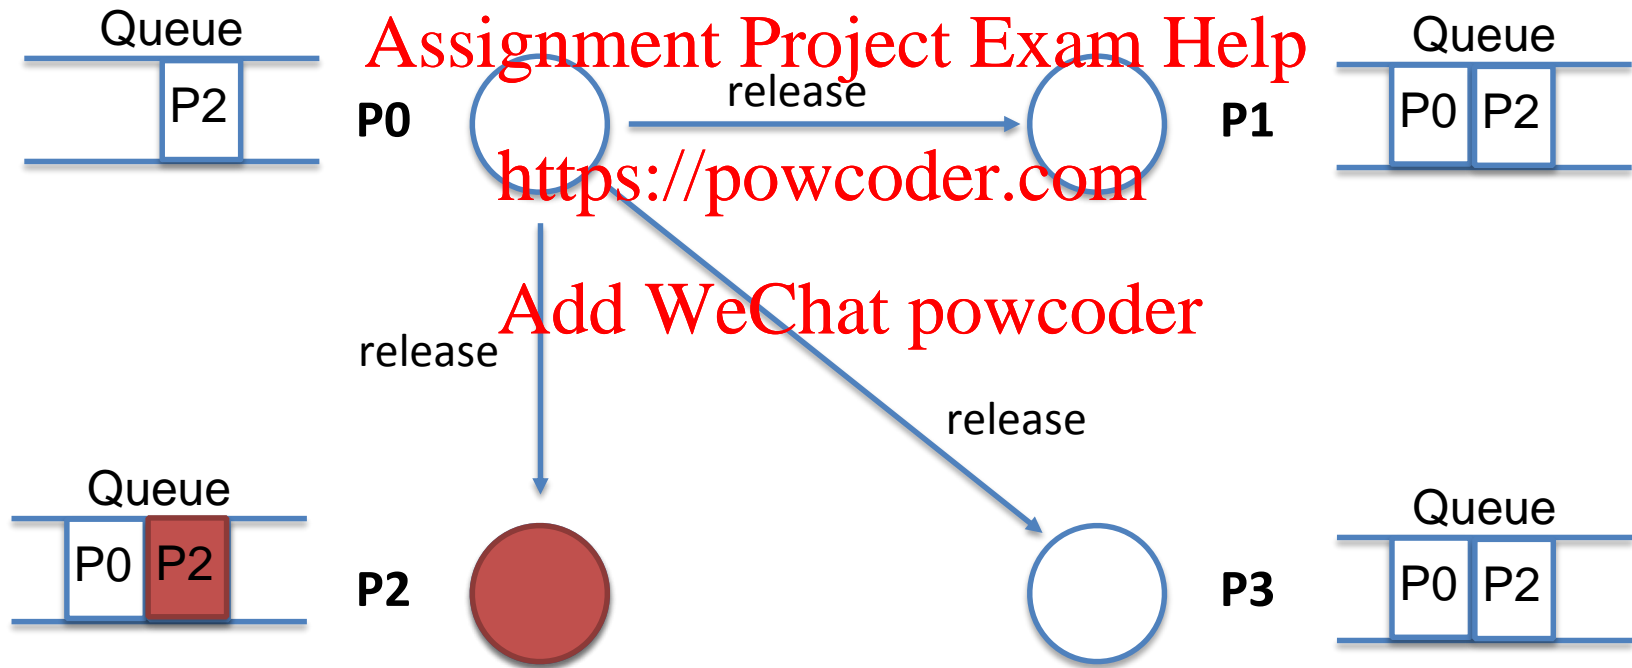
# Example

Using process id to break ties, process P0's request will end up on top of the queue, over process P2's.



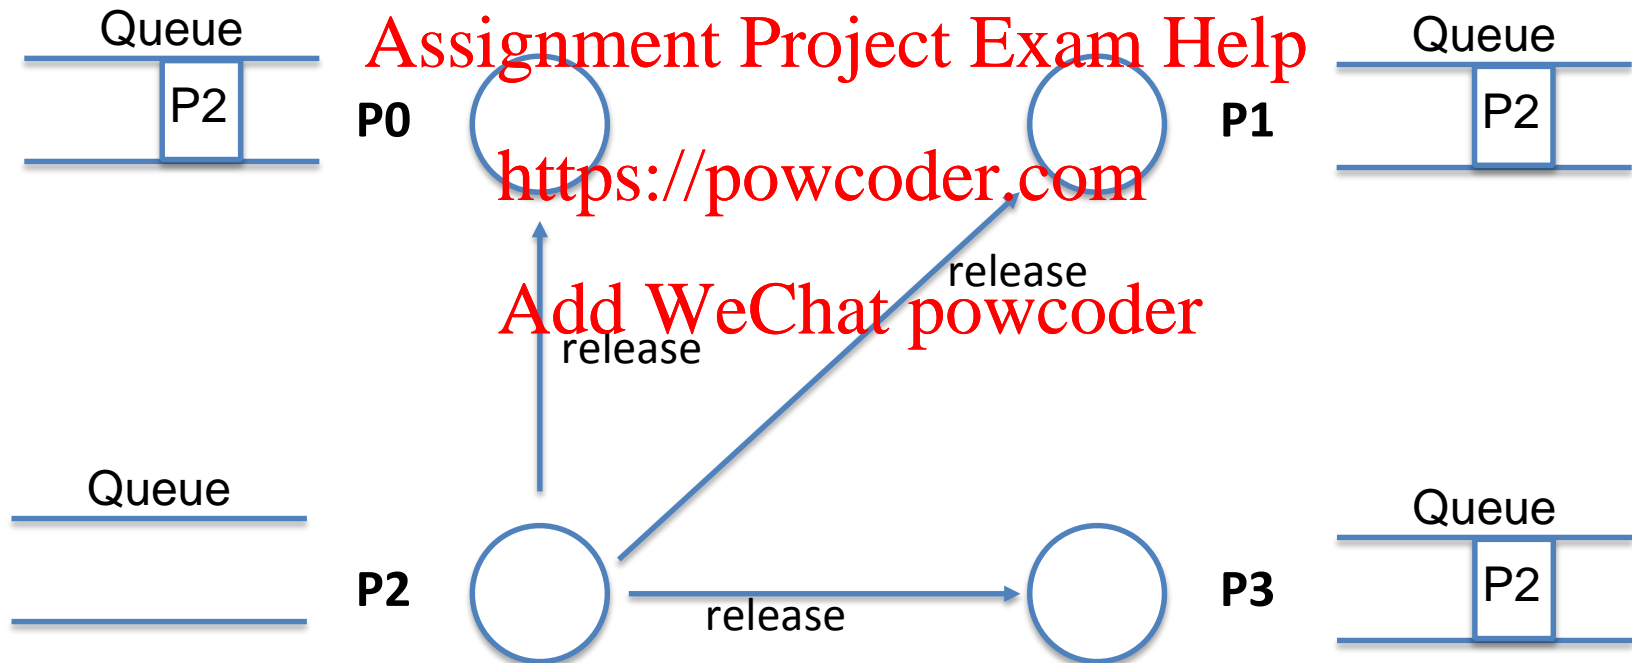
# Example

Process P0 is done and broadcasts a release message.



# Example

When process P2 is done with the shared resource, it also broadcasts the release message.



# Complexity of Lamport's Solution

- This algorithm creates  $3(N - 1)$  messages per request:
  - $(N - 1)$  total number of requests
  - $(N - 1)$  total number of replies
  - $(N - 1)$  total number of releases
- Can we reduce the number of messages while maintaining the same guarantees?

# Quorums

- What if there are **failures**?
  - Mostly network failures
  - No failure of the resource holder
- Do we need to communicate with **ALL** processes? <https://powcoder.com>

**Add WeChat powcoder**

- A **quorum** is the minimum number of **votes** that a process has to obtain in order to be allowed access to the shared resource.

# Quorums

- Any two requests should have a common process to act as an **arbiter**.
- $V_i$  is called a **quorum**.
- Let process  $p_i$  ( $p_j$ ) request permission from  $V_i$  ( $V_j$ ), then

<https://powcoder.com>

Add WeChat powcoder



# Quorums

- Given  $N$  processes:  $|V_i| > N/2$ , ie,

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- In general, majority, ie  $\lceil (N/2) \rceil$ .  
[Gifford 79]

# Basic MUTEX Protocol

- Requesting CS

- process requests CS by sending *request* message to processes in its quorum
- If a process receives a *request* it sends back *reply* unless it granted permission to other process; in which case the request is queued—similar to granting a lock

- Entering CS

<https://powcoder.com>

- process may enter CS when it receives *replies* from all processes in its quorum

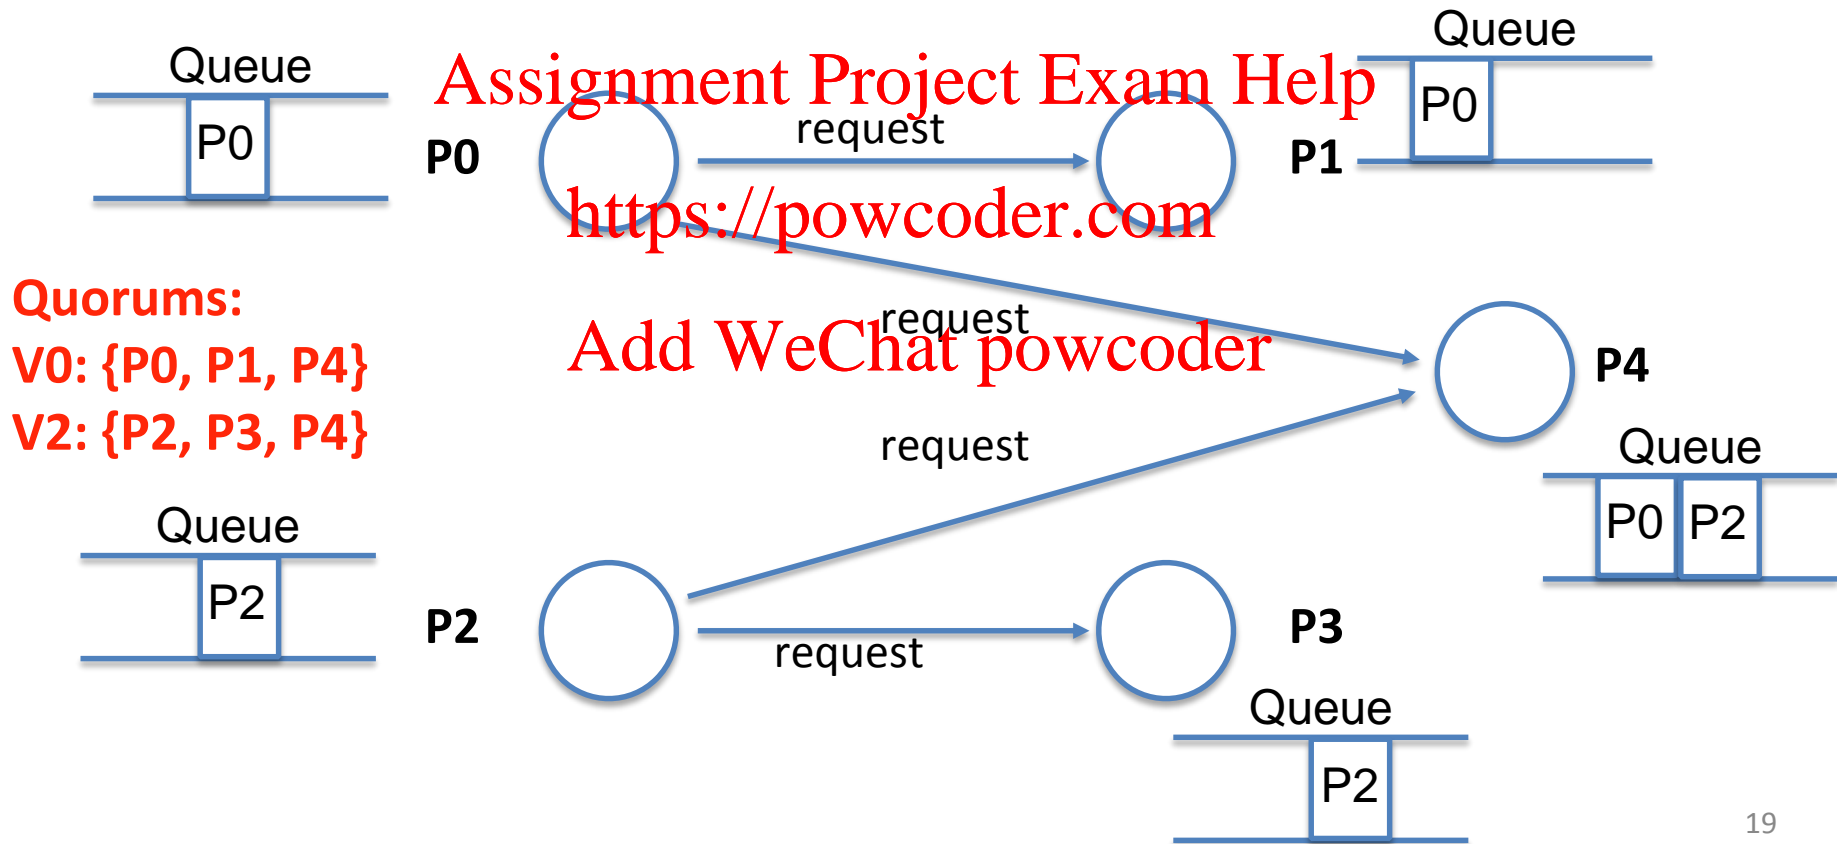
Add WeChat powcoder

- Releasing CS

- after exiting CS process sends *release* to every process in its quorum—release lock
- when a process gets *release* it sends *reply* to another request in its queue

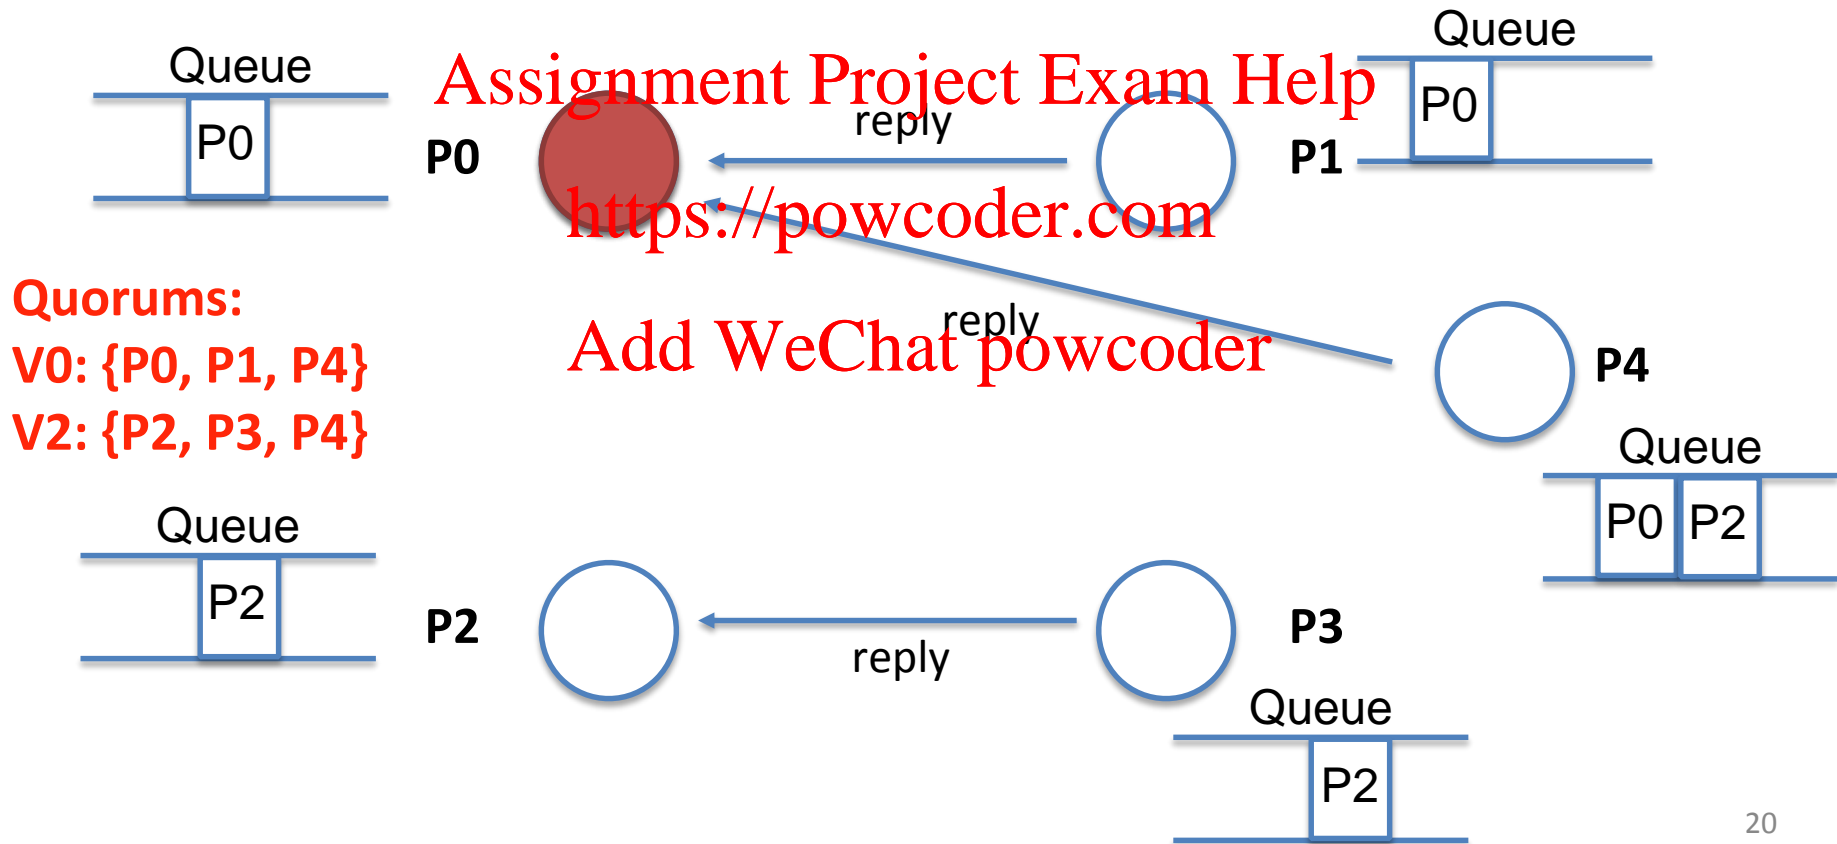
# Example

Assume that process P0 and process P2 need to access the shared resource at the **same time**.



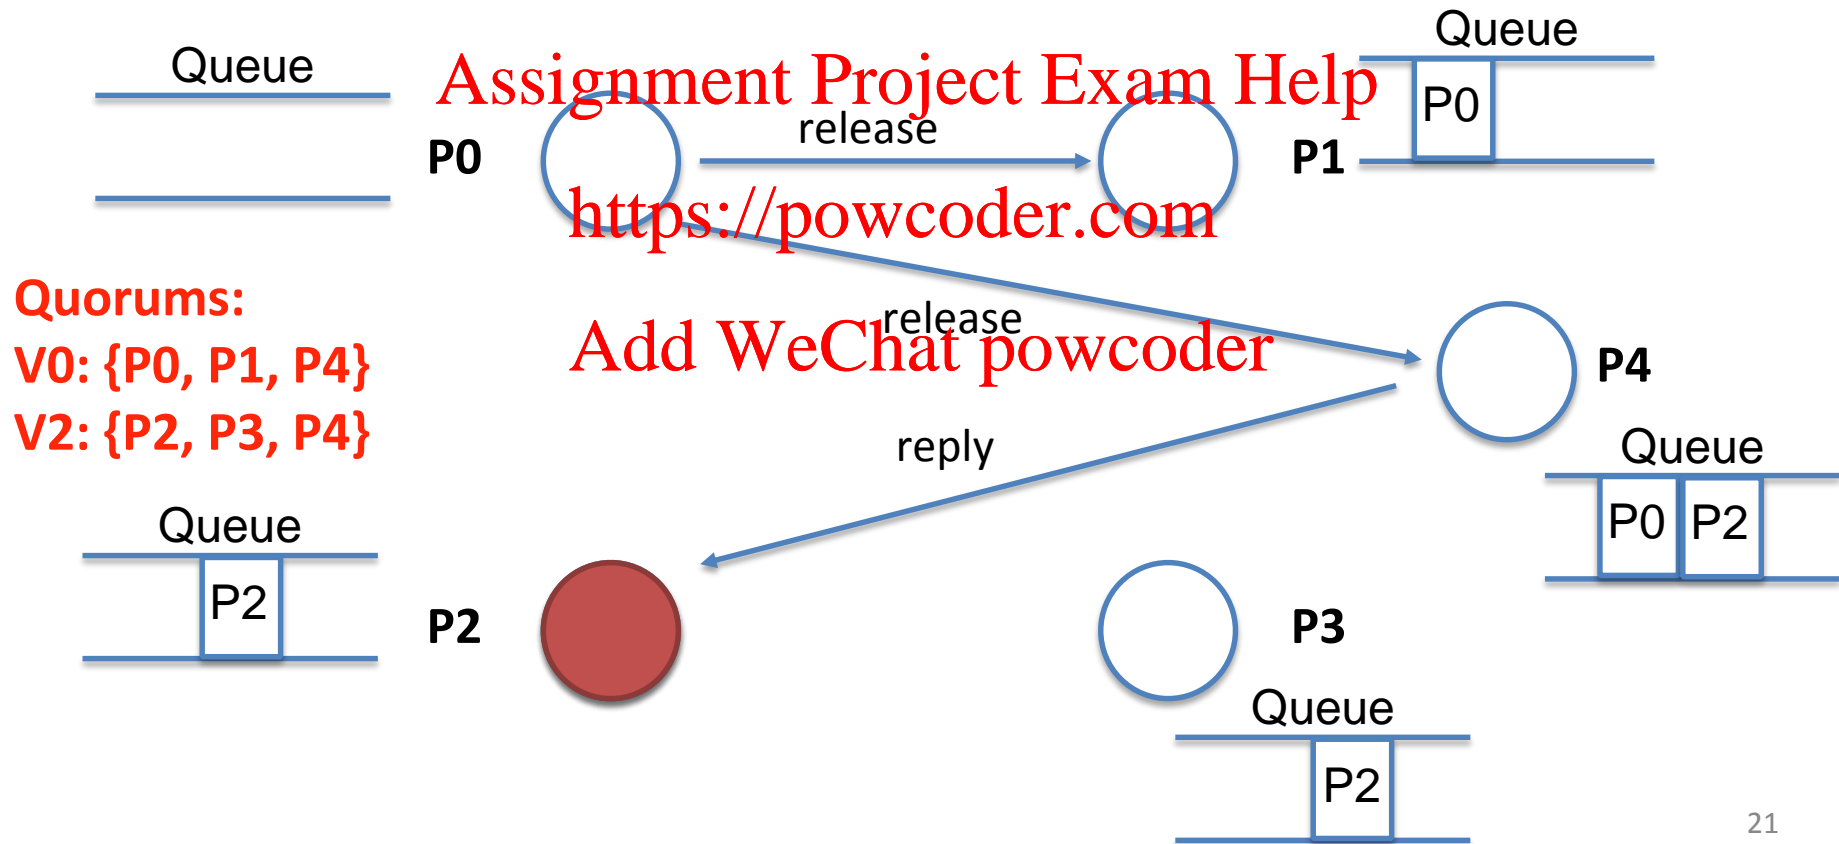
# Example

Process P4 is the arbitrator since it's common in both V0 and V2 quorums.



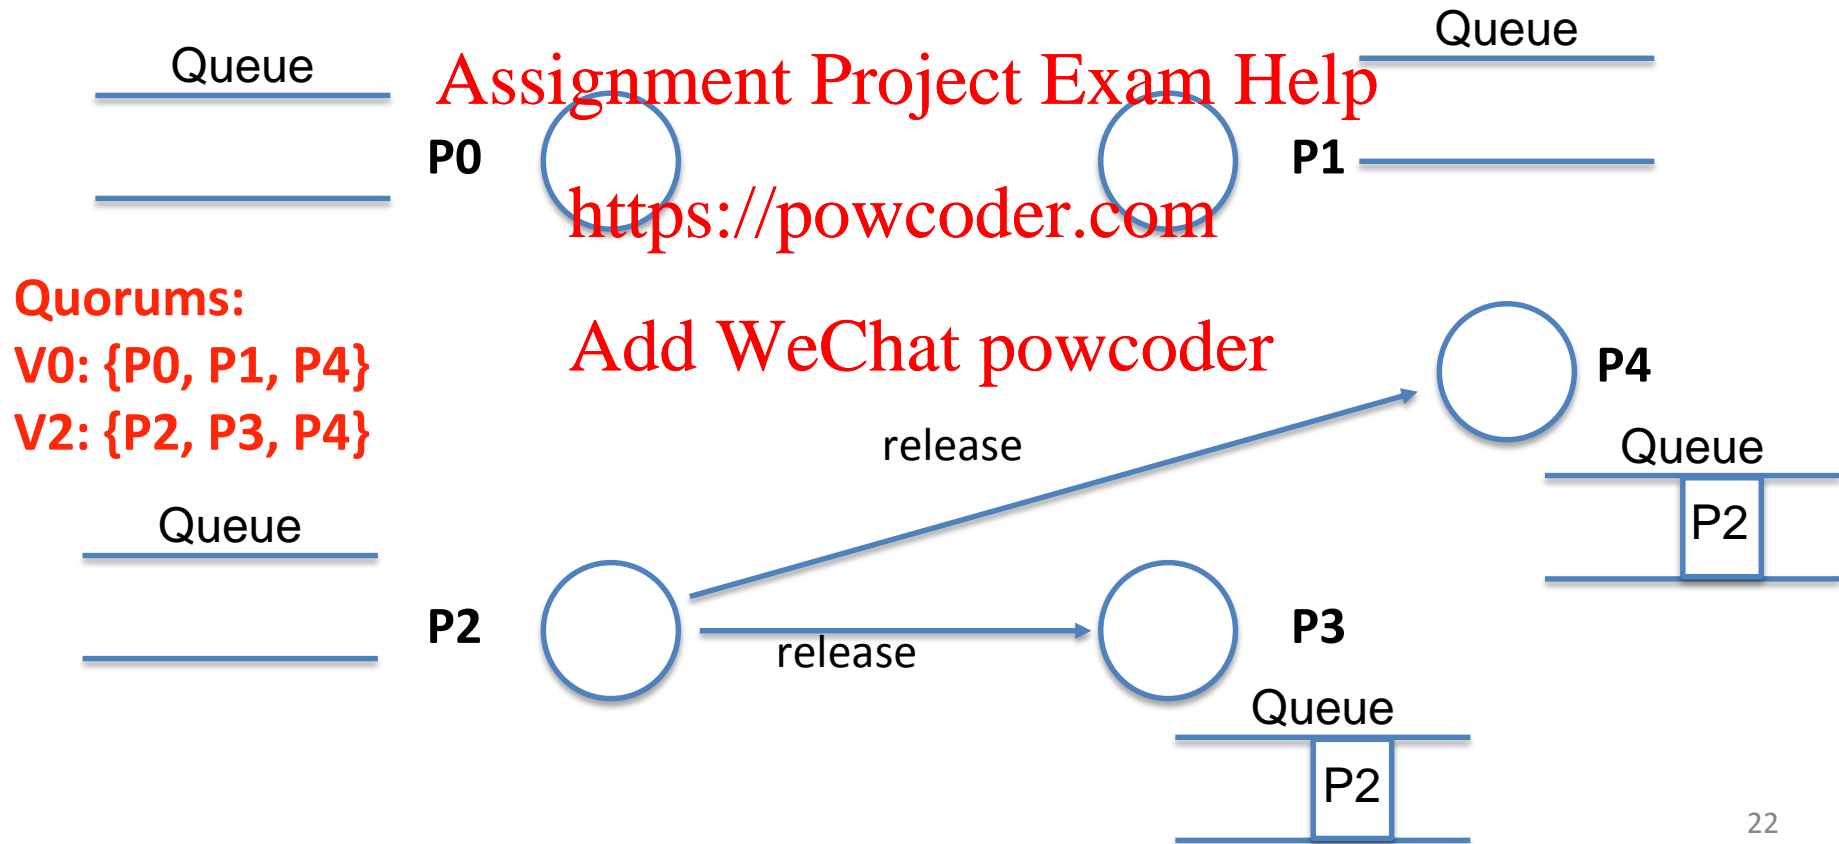
# Example

Process P0 sends the release message to its quorum.



# Example

Process P2 finishes and broadcasts to its quorum the release message.

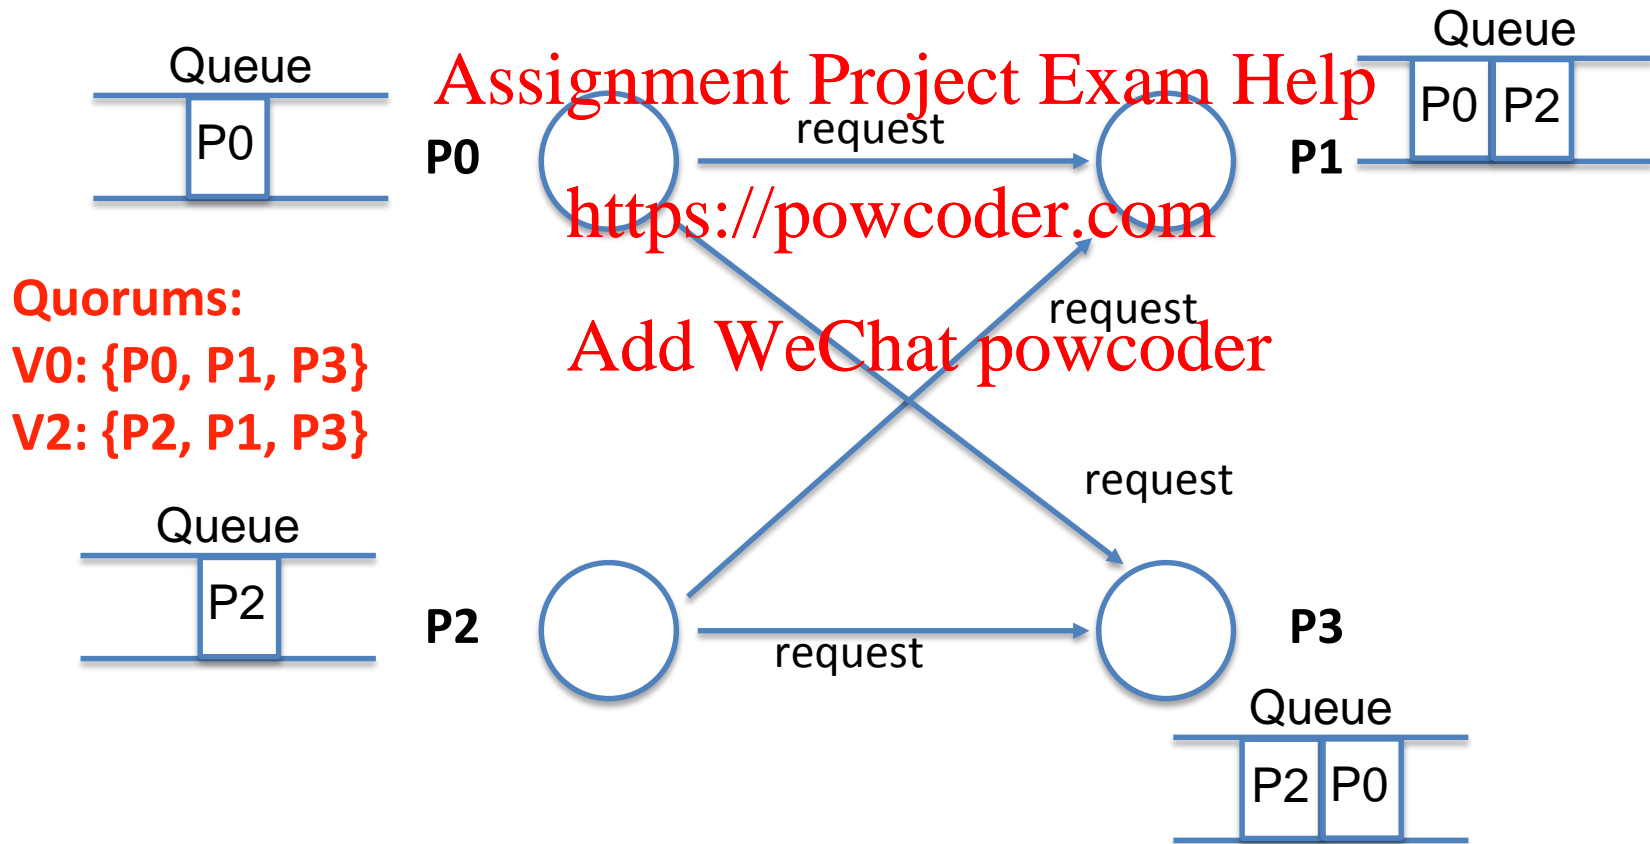


# Deadlock

- Unfortunately, the quorums approach can lead to **Deadlocks**.
- Deadlocks occur even with a **single resource**.
  - Not the standard database or operating system multi-object deadlock.

# Deadlock Example

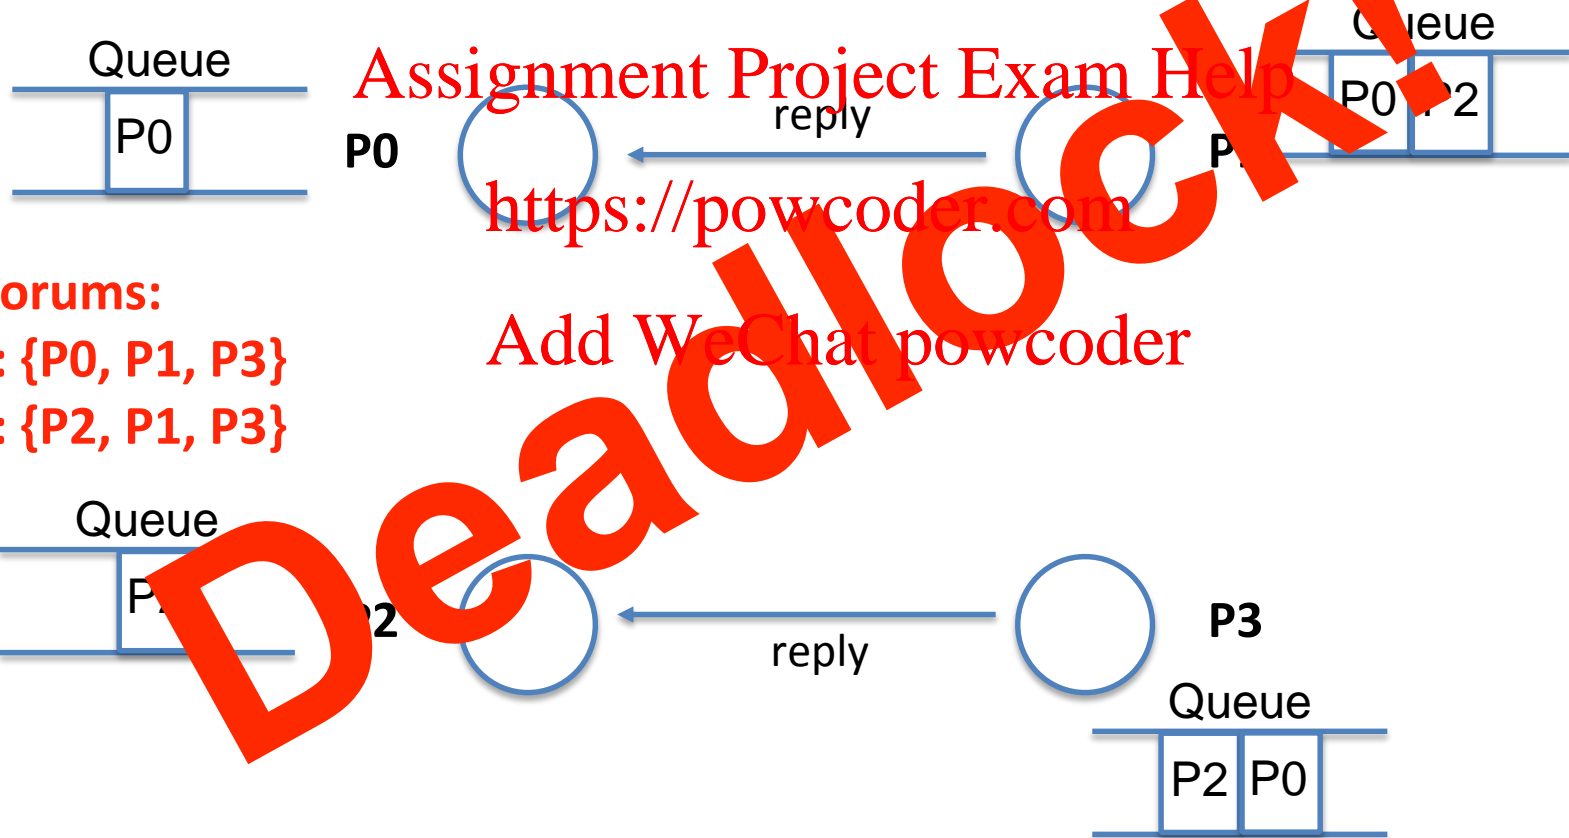
Assume that process P0 and process P2 need to access the shared resource at the **same time**.





# Deadlock Example


P1 and P3 don't agree on which process should get the resource.



# Site Failure Example

How many sites can we afford to go down?

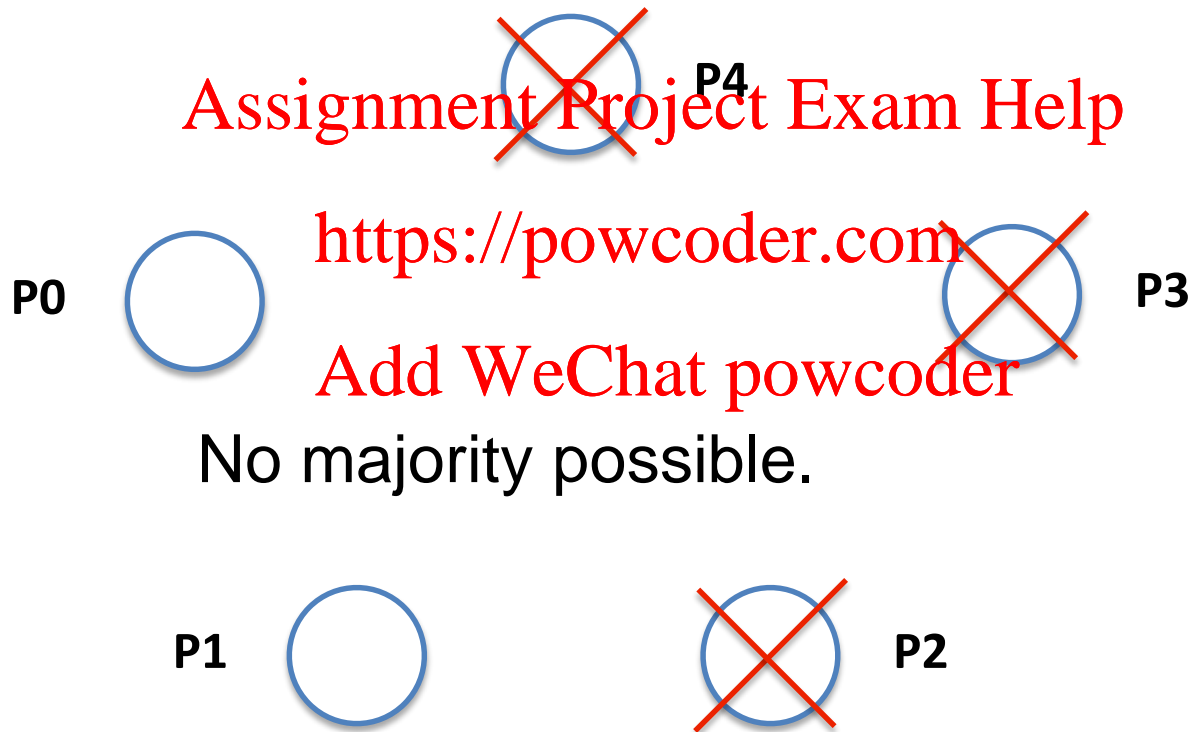
We can achieve quorum!



# Site Failure Example

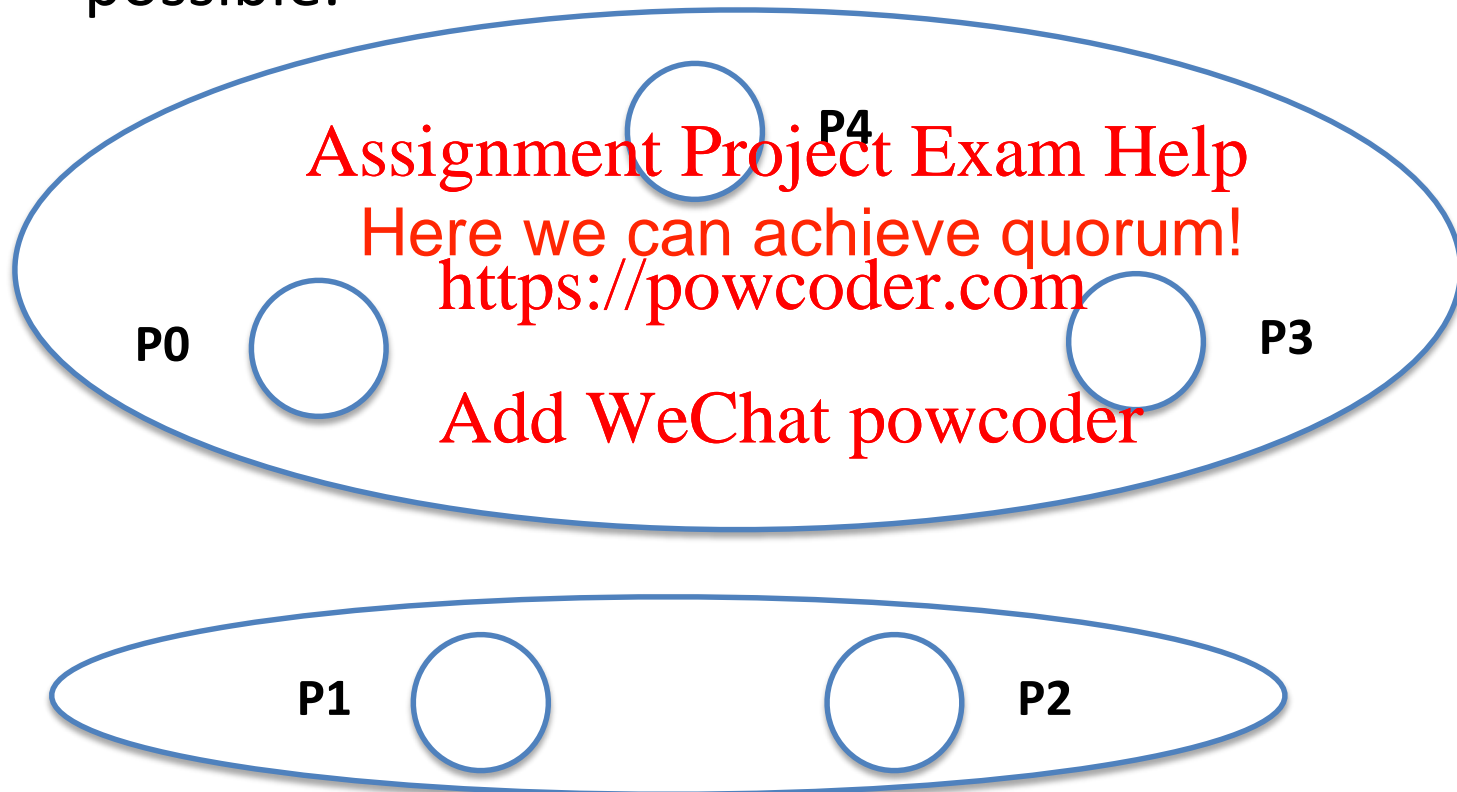
How many sites can we afford to go down?

We can **not** achieve quorum!



# Network failure: Partitioning

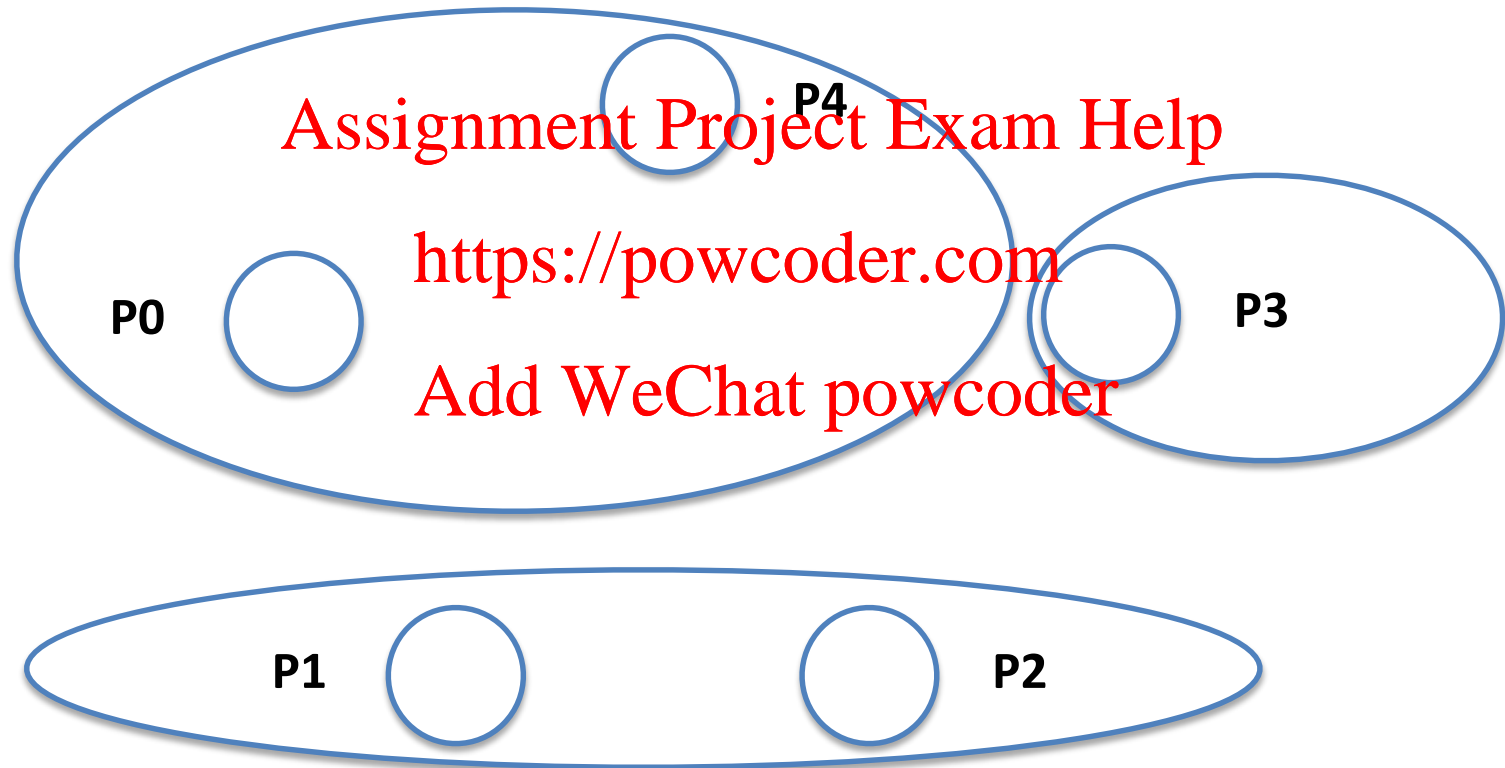
If sites are unreachable, quorum might still be possible.



This partition won't be able to get the resource.

# Network failure: Partitioning

If sites are unreachable, quorum might still be possible.



**No** partition can now get the resource.

# Studying the Sizes of Quorums


- Can we have quorums of sizes **LESS** than majority?
- Yes:
  - Using logical structures

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Quorum Sizes on Logical Grid Structure



Quorum size is  $O(\sqrt{N})$